

Contextual effects on online pragmatic inferences.

JK; JL; MC

Affiliation

## Abstract

Where the veracity of a statement is in question, listeners show a bias towards interpreting speaker disfluency as a sign of dishonesty. This bias is not limited to post-hoc judgements, but can also be found during online speech processing. The present study investigates whether listeners are influenced by contextual information about the potential causes of speaker disfluency. If listeners make inferences about the cause of a disfluency, then a plausible speaker-distraction may attenuate the inclination to interpret an utterance as dishonest.

Participants listened to a speaker describe the locations of treasure (“The treasure is behind [the]/[thee—uh—] <referent>.”), while viewing scenes comprising the referent and a distractor. They were told that not all utterances would be honest, and their task was to click on the suspected location of the treasure. In line with previous work, participants were more likely to click on the distractor when the description was disfluent, and this effect corresponded to an early fixation bias, demonstrating the online nature of the pragmatic judgement.

The present study, however, also manipulated the presence of a plausible external cause of speaker disfluency. To accomplish this, participants were told that all utterances were produced outside on a busy street, and all items were played over low-level ambient street noise.

When there was a plausible external cause of the speaker disfluency (a distracting noise, in the form of a car-horn), participants were more likely to initially fixate on the referent, only later fixating on and selecting the distractor. One account of these findings would suggest that participants made early inferences about the contextual causes of disfluencies, which were eventually overridden by a dishonesty bias for disfluent utterances.

Contextual effects on online pragmatic inferences.

### Background.

During conversation, rarely is an utterance or its delivery pre-planned. Instead, everyday speech is for the most part spontaneous and thus often disfluent, containing pauses, “um”s, “uh”s, repetitions, revisions, and mispronunciations. Naturally occurring speech has a rate of approximately 6 to 10 disfluencies per 100 words [Bortfeld et al., 2001; Fox Tree, 1995]. The disfluent nature of speech is just one of many variable aspects of *how* an utterance might be presented, and listeners must be able to cope with this variability to successfully interpret a speaker.

These paralinguistic features of conversation are not merely incidental—often an utterance’s manner of delivery is apposite to the message being conveyed, from the prosodic contour [Fernald and Mazzie, 1991], to an accompanying gesture [Alibali et al., 2001]. In this way, numerous paralinguistic ‘cues’ about the content of a speaker’s message are readily available to listeners, and research has shown that listeners can, and do, exploit these phenomena [Corley et al., 2007; Barr, 2001; Frazier et al., 2006]. However, the way in which listeners interpret paralinguistic cues is not clear—do they simply rely on pre-learned direct associations of specific paralinguistic features with certain events in language, or do they use social and contextual information to draw inferences about why an utterance was produced in the way that it was? The latter suggestion would amount to a form a modelling the speaker’s intentions, thoughts and perspective.

Disfluencies have been shown to vary with a speaker’s planning processes, with rates of disfluency in speech increasing along with a speaker’s cognitive load[Bortfeld et al., 2001]. Research shows that speakers are more disfluent when utterance planning involves low-frequency words [Beattie, 1979], less-preferred syntactic structures [Cook et al., 2009], discourse-new expressions [Arnold et al., 2003], and more choice of expressive alternatives [Schachter et al., 1991]. These may in turn provide the listnere with information about the speaker: listeners appear to associate speaker-disfluency with aspects of language which are perceived as more difficult to plan, displaying biases towards less predictable continuations [Corley et al., 2007], discourse-new referents [Barr, 2001; Arnold et al., 2004], and unfamiliar objects [Arnold et al., 2007].

The sensitivity which listeners display to the manner of an utterance’s delivery facilitates the processes involved in comprehension of the linguistic message by generating expectations about upcoming speech. However, listeners’ biases to disfluencies are not restricted to their predictions about the linguistic message, and also extend to global

pragmatic judgments made about an utterance, such as estimates of the speaker's metacognitive state. Utterances which are hesitant, disfluent, or with rising intonation are judged as less confident [Smith and Clark, 1993; Brennan and Williams, 1995]. This effect has shown to be present in various contexts—in both the speaker's own and a listener's assessment of certainty [Swerts and Krahmer, 2005], and for both adult and child listeners [Krahmer and Swerts, 2005].

When a speaker lies, however, we might perceive their utterance planning as more effortful, in order to convincingly convey a statement which they do not regard to be true. Speaker-disfluency in particular, and especially pauses, appears to be frequently associated with an interpretation of dishonesty, in judgments made about speakers themselves as well as judgments made about other speakers [Zuckerman et al., 1981]. Research has shown that accuracy in discriminating truth from lies increases when presented with audio or audiovisual information rather than visual information [Bond and DePaulo, 2006], supporting the salience of vocal cues in recognising deception.

The effect of disfluency on listeners' pragmatic judgments about an utterance have, until recently, used offline measures of judgement (post-hoc assessment of a speaker's confidence or honesty). However, recent research has found that the disfluency-dishonesty bias is apparent from the early stages of reference comprehension. In a recent study by Loy et al. [2016], subjects engaged in what was posed as a 'lie detection game', in which listeners were tasked to make implicit judgments about whether a speaker is being deceitful or not in indicating the location of a reward. Utterances were presented as either fluent or disfluent ("The treasure is behind [the]/[thee—uh—] <referent>."). In a visual world eye- and mouse-tracking paradigm, listeners signalled their assessment of a speaker's honesty by clicking on one of two objects (the referent and a distractor) - interpreting an utterance as honest would result in a click on the referent, and as dishonest a click on the distractor. In this way, Loy et al. [2016] investigated how listeners' pragmatic inferences about a speaker's honesty during online processing of speech was influenced by speaker-(dis)fluency.

Loy et al. [2016] found that, as with the influence of disfluencies on listeners' predictions about semantic content [Arnold et al., 2004, 2007; Barr, 2001], manner of delivery's influence on listeners' pragmatic hypotheses emerges rapidly during the initial stages of reference comprehension. Subjects' eye- and mouse-movements displayed a bias toward the referred-to object for fluent utterances, and to the distractor object for disfluent utterances [Loy et al., 2016]. Furthermore, these biases were found for utterances with both utterance-initial and utterance-medial disfluencies, despite the general belief that

utterance-initial disfluencies are associated more with global speech-planning such as that involved in deception. This dishonesty-bias listeners display to speaker-disfluency shows evidence of an ability to integrate information about utterance delivery and alter the global interpretation of an utterance during the moment-to-moment processing of speech.

The aforementioned studies give strong evidence for listeners having biases towards speaker-disfluency in terms of both semantic content [Arnold et al., 2004; Barr, 2001] and pragmatic judgments [Loy et al., 2016], and for the emergence of these biases during the initial stages of the comprehension process. The mechanism by which disfluency influences listeners interpretations, however, remains less clear. The disfluency-biases discussed above can all be explained in terms of a low-level account in which listeners’ interpretations of disfluencies result from a prelearned stochastic model of co-occurrences of disfluencies with certain properties of language use. Listeners’ bias towards associating presence of disfluency with discourse-new information could be due to distributional learning of disfluencies more often occurring when speakers refer to something new than to something familiar [Arnold et al., 2004; Barr, 2001]. Likewise the disfluency-dishonesty bias in Loy et al. [2016] could result from a prelearned association (or belief of an association)? between disfluency and deception.

An alternative account suggests that speaker-disfluency is a form of ‘collateral communication’ [Clark, 1996]. As opposed to holding direct associations of disfluencies with language use, this view suggests that listeners interpret disfluency as communicative signals which manage discourse, and generate expectations by drawing inferences about its cause—e.g. about what might cause increased cognitive load for the speaker. This ‘disfluency-attribution’ view is paired with the suggestion that speakers’ production of fillers is akin to the use of conventional words, carrying a fundamental meaning of delay to which the listener can then attribute a cause [Clark and Fox Tree, 2002]. For the disfluency-deception bias, this account suggests that listeners interpreted disfluency by drawing an inference that e.g. the speaker’s decision to lie resulted in more effortful utterance planning, thus causing disfluency.

Understanding the mechanisms underlying listeners’ processing of disfluencies has wider implications for language comprehension in general. If listeners inferentially attribute cause to disfluency during online speech processing, this suggests that they rapidly access information about what might cause delay for that particular speaker in that particular context. As such, this account implies that listeners engage in a form of modelling of the speaker. In comprehension, accounting for a speaker’s perspective has been argued to be

costly in effort [Lin et al., 2010], as it requires continually accessing speaker- and context-specific information. Instead, many have suggested that listeners rely on less cognitively demanding, ‘egocentric’ tactics, in which language is interpreted relevant to a listener’s own perspective (or a minimally adapted version), and taking a speaker’s perspective occurs only when entirely necessary [Barr, 2004, 2014; Pickering and Garrod, 2004; Kronmüller and Barr, 2007]. If listeners’ interpretations of disfluencies are moderated by speaker- and context-specific information, this suggests that the effect of disfluencies on comprehension cannot be the result of prelearned associations alone. Instead, it would suggest that comprehension of disfluent speech involves taking into account what might lead a particular speaker to produce an utterance in a particular way, and as such, implies a form of perspective-taking.

Recently, several studies have provided evidence for disfluency-comprehension being context-dependent by examining how listeners’ biases to disfluencies change when presented with alternative plausible causes of speaker-disfluency. Arnold et al. [2007] showed that a bias towards unfamiliar, hard-to-describe objects following a disfluency is moderated by listeners’ beliefs about that speaker. The results of both a gating task and an eye-tracking task showed that a) the discourse-new bias found in Arnold et al. [2004] could be extended to an unfamiliarity-bias, and b) this bias was reduced when the listener believed the speaker to have object-agnosia (inability to recognize objects). Similarly, Barr and Seyfeddinipur [2010] found that the discourse-new bias to disfluencies was dependent upon what was new for the speaker independent of what was new for the listener. Heller et al. [2015] suggest that listeners use context-specific inferences about what objects may cause production delay—they found that listeners’ bias towards unfamiliar objects adapted to novel contexts with artificial lexicons in which all objects were, to varying degrees, ‘unfamiliar’.

Whilst these studies show that listeners’ disfluency-biases display a flexibility to speaker- and context-specific information, the extent to which listeners are engaging in inferential processing about causes of disfluency is disputable. For listeners’ perceptions about the cause of disfluency, belief that the speaker is object-agnosic provides an utterance-global explanation for speaker-disfluency (familiar and unfamiliar objects being equally difficult to describe for an object-agnosic speaker). In this respect, whilst the integration of information evident in Arnold et al. [2007], Barr and Seyfeddinipur [2010], and Heller et al. [2015] clearly requires some form of inferencing about the speaker or context, it is possible that listeners only engaged in the inferential process prior to speech onset. This could equate to essentially readjusting prior associations of disfluency and unfamiliar objects,

rather than repeatedly making inferential judgments about cause of disfluency as speech unfolds.

To investigate whether listeners access context-specific information to engage in disfluency-attribution *during* the online processing of speech, we ask whether listeners' disfluency-biases are flexible with regard to the presence of an utterance-local cause of speaker-disfluency—i.e. an explanation of disfluency which only presents itself *during* the time-course of an utterance. Specifically, we look at whether listeners' pragmatic biases to disfluencies signalling deception [Loy et al., 2016] is moderated by the availability of 'speaker-distraction' as a cause of disfluency.

The current study develops Loy et al. [2016]'s visual world paradigm, with the addition of manipulating the presence of a disruptive (and so plausibly distracting) noise in the speaker's environment. Thus the study followed a 2 (fluent/disfluent) x 2 (distraction absent/present) design. To accomplish this, subjects participated in the same 'lie-detection task', but this time utterances were presented to listeners under the guise of having been recorded outside in a busy street. Low-level ambient noise was present behind every utterance, and the potentially distracting noise took the form of a car-horn. As in Loy et al. [2016], we look at eye- and mouse- movements, to study the time course of listeners' pragmatic judgments about the honesty of an utterance. We predict that, for utterances in which there is no car-horn present, results will replicate the findings of Loy et al. [2016]—we expect to see a bias towards the referent (and an implicit interpretation of honesty) for fluent utterances, and towards the distractor for disfluent utterances. If listeners draw inferences about the cause of speaker-disfluency during the moment-to-moment processing of speech, we would expect to see an attenuation of the bias towards the distractor object for disfluent utterances with the car-horn distraction present.

## Method.

### Materials.

Visual stimuli consisting of 120 black and white line drawings, taken from Snodgrass and Vanderwart [1980], were presented to participants in pairs across 60 trials (20 experimental, 40 fillers). Each trial presented the *referent* (the object which the speaker identified as having the treasure behind for that trial), and a *distractor*, which was chosen at random without replacement from a set of 60 objects. To control for the effect of the bias towards interpreting disfluency to difficulty of description [Arnold et al., 2007], critical referents and distractors were matched for familiarity (H value  $\geq 3.0$ ) and ease-of-naming

(< 1.0). Object pairings with the same phonetic onset were avoided.

As in Loy et al. [2016] each referent was associated with a recording of a speaker naming the depicted object as the location of the treasure, and as before, critical utterances were either fluent or disfluent. The present study, however, manipulated not only the manner of delivery, but also the presence of a plausible cause of speaker distraction. Thus audio files were constructed such that the critical referents could be heard in four conditions in a 2 (fluent/disfluent) x 2 (distraction absent/present) design, see below. The insertion of a prolonged article followed by a filled pause into each fluent utterance formed the disfluent variants, ensuring that fluent/disfluent counterparts were identical save for disfluency manipulation. Where Loy et al. [2016] studied both utterance-initial and utterance-medial disfluencies, we opted for an utterance-medial position as it was considered more believable as being caused by speaker distraction when paired with a environmental cue. Meanwhile, to create a plausible cause of speaker-distraction, a 520ms car-horn sound effect was presented prior to the onset of the referent noun (-1100ms for disfluent utterances, -600ms for fluent).

1. **fluent, distraction absent:** The treasure is behind the <referent>.
2. **disfluent, distraction absent:** The treasure is behind thee, uh ... <referent>.
3. **fluent, distraction present:** The treasure is behind the <referent>.  
horn
4. **disfluent, distraction present:** The treasure is behind thee, uh ...<referent>.  
horn

To situate the car-horn sound effect in a context, all recordings were presented overlaid upon a clip of ambient traffic and street noise. These clips were unique to each referent, so that the presentation of each critical referent varied only by 1) disfluency manipulation and 2) presence of car-horn. This ensured that participants' behaviour was a reaction to the different combinations of these elements. All recordings, ambient noise and sound-effects were normalised and resampled to create 48 KHz stereo .wav files.

The 20 critical items/referents were counterbalanced across four lists, each with 10 fluent and 10 disfluent utterances, and each containing 10 instances of the car-horn (5 of which were disfluent). Lists also contained 40 filler utterances, half of which were fluent, and half of which contained either a form of disfluency or a discourse manipulation (see Table 1). Additionally, 20 of these filler items (10 fluent, 10 disfluent/discourse manipulation) contained various novel noises which could be interpreted as distracting for a speaker (see Table 2). These filler distractions varied in position relative to the referent noun onset.



**Cover story.**

Central to the idea of a distraction-caused-disfluency is the requirement that participants believe that the utterances were produced naturally and in a noisy environment. Participants were told that the recordings were made by a participant of a previous experiment which was conducted on the side of a busy street. To corroborate this, during the initial explanation of the experiment, participants were shown three videos claimed to be examples of the speaker producing the utterances. These were presented alongside the images about which the speaker spoke (see Figure 1). Speaker utterances were depicted in the first video as honest & fluent, dishonest & disfluent in the second, and, in the third, honest & disfluent but distracted (speaker glances sideways toward a car following the car-horn sound effect).

To ensure that analysis could be run only on data from participants for whom the cover story held, a post-test questionnaire assessed whether participants thought it believable that the utterances were produced naturally and in the suggested context. This was double-checked by verbal questioning after participants were debriefed about the true nature of the experiment, by asking whether it occurred to them during the experiment that the recordings might not have been produced outside in the street.

**Procedure.**

Stimuli were displayed on a 21" CRT monitor, placed 850mm from an Eyelink 1000 Tower-mounted eye-tracker which tracked eye-movements at 500Hz (right-eye only). Audio was presented in stereo from speakers either side of the monitor. Mouse coordinates were sampled at 50Hz. The experiment was presented using OpenSesame version 3.0 [Mathôt et al., 2012].

Participants were told they would see two objects and hear a speaker identify one of them as hiding the treasure. They were told that the objective of the game was to accrue treasure by successfully locating the treasure, the location of which the speaker would lie half the time about. Participants' instructions were to click on the object which *they believed* the treasure to be behind. Once participants had read the instructions, the eye-tracker was calibrated. Calibration was monitored throughout, with the experiment being paused for recalibration if the experimenter deemed it necessary. To maintain accuracy, between each trial participants underwent a drift correction using a central fixation point, which changed from grey to red (for 500ms) upon successful fixation, signifying the start of the subsequent trial. Each trial began with the presentation of the two images (referent and distractor) in

positions vertically central and horizontally to the left and right of centre (referents were presented equally often on each side). Simultaneous with presentation of the visual display was the onset of the ambient traffic audio file. After 1500ms, the mouse pointer appeared centrally, and the playback of the utterance began. Upon clicking an object, the stimuli disappeared and were replaced by the grey fixation dot, signifying progression to the subsequent trial. Each trial had an automatic time-out of 5000ms from utterance onset, in case participants did not click on an object.

To maintain motivation throughout the study, participants were told that there were a number of “hidden bonus rounds” which offered more treasure. Of the filler trials, 25% did not immediately progress to the subsequent trial, instead giving a positive feedback message (regardless of which object was clicked) stating successful completion of a bonus round. Furthermore, participants were told that top scorers of the game would be able to enter their name on a highscore table, which was shown at the beginning of the experiment. Participants completed five practice trials (one of which was presented as a bonus round) prior to the main experiment. Eye-movements, mouse co-ordinates and object-clicked (referent/distractor) were recorded for each experimental trial.

## Results.

### Exclusion criteria.

For a final sample of 24, we were required to run the experiment with 37 participants, recruited from the University of Edinburgh community, who participated in return for £4 remuneration. The data from 12 participants was excluded on the basis that they indicated either in the questionnaire (ten) or verbally (two) that they did not believe the cover story. Additionally, the study reported that one participant for whom the deception held had previously participated in [Loy et al. \[2016\]](#), and thus they were also excluded. Analysis was run on the data from the remaining 24 participants.

### Analysis.

Analysis was carried out in R version 3.3.0 [[R Core Team, 2016](#)], using the lme4 package [[Bates et al., 2015](#)]. Trials in which participants did not click on either the referent or distractor (0.01%) were excluded from all analyses. Object-clicked (referent/distractor) was modelled using mixed effect logistic regression, testing for main effects of manner of delivery, presence of plausible speaker distraction, and their interaction as fixed effects, and

including random slopes for these same variables, plus random intercepts and slopes by both subject and referent.

Eye-tracking fixation data was averaged into 20ms bins (of 10 samples) prior to analysis, and coded in terms of the area of interest (AOI), which were referent; distractor and neither image. The proportion of fixations to either object from the total sum of fixations was computed for each time bin. Mouse-tracking analysis only looked at the  $X$ -coordinates of the cursor. Distance traveled by the cursor (pixels) was computed for each 20ms bin, coded for direction relative to either object. The cumulative mouse distance traveled toward either object in each trial was calculated by summing the distance travelled up to each time bin, and proportion of movement was calculated by taking the distance travelled toward an object, divided by the total distance travelled in that trial in either direction. Trials for which the total mouse distance travelled post referent-onset was less than 1/3 of the distance from the screen centre to the near edge of an object were excluded (0.03% of trials), on the basis that it was indicative of participants' having decided prior to referent-onset upon which object to click. Movements beyond the outer edge of either object were excluded on the basis that it indicated 'overshooting' with the cursor.

Model analysis for both eye- and mouse-movements were conducted over a time window from referent onset to 800ms post-onset. This window coincides with [Loy et al. \[2016\]](#), identified by evidence that reference is established in eye-movements around the 400-800ms post noun-onset mark [[Eberhard et al., 1995](#)], and covering the duration of the longest critical referent (776ms). Models were estimated using empirical logit regression [[Barr, 2008](#)] with main effects of time (scaled), object (referent/distractor), delivery (fluent/disfluent) and speaker distraction (absent/present), and interactions between them all, along with by-subject and by-item random intercepts and random slopes for time.

### **Final object-click.**

Responses show the same overall tendency to interpret an utterance as truthful as was found in [Loy et al. \[2016\]](#), with 57% of trials resulting in a click on the referent and only 43% on the distractor. Table 4 shows the percentage of mouse-clicks on each object by condition. Participants were less likely to click on the referent following a disfluent utterance than a fluent one ( $\beta = -2.24, SE = 0.67, p < .001$ ). This is in keeping with literature [[Brennan and Williams, 1995](#); [Swerts and Krahmer, 2005](#)] and with the results of [Loy et al. \[2016\]](#) - manner of delivery influences participants' global interpretation of the speaker's truthfulness. The presence of a plausible speaker distraction was not found to affect responses, neither was the

interaction between delivery and distraction. The bias toward interpreting disfluency as a sign of dishonesty appeared to be explicit for 19/24 participants, as indicated in the post-test questionnaire.

### Eye-movements.

The time-course (referent onset to 2000ms post-onset) of fixations to referent and distractor can be seen for the fluent conditions (Fig. 2) and the disfluent conditions (Fig. 3). Visualisation of the time-course of fixations show that for fluent utterances, participants displayed an early fixation bias towards the referent, which began to present itself prior to mean critical noun offset. For disfluent utterances, the picture is more complex—whilst a fixation bias towards the distractor was present, it appeared later on in the trial, not during the utterance of the critical noun as it did in Loy et al. [2016]. Although the same is true of utterances in which there was an available cause of speaker-disfluency, in these cases, participants appeared to fixate more to the referent first, before the later bias towards the distractor.

Model analysis of the window from referent-onset to 800ms post-onset confirmed what is evident from visualising the time-course of fixations. When presented with a fluent utterance, participants' fixations to the distractor decreased over time in favour of the referent ( $\beta = -0.63, SE = 0.03, t = -18.51.$ ). In contrast, fixations to the distractor over the referent increased over time when the delivery was disfluent ( $\beta = 0.59, SE = 0.05, t = 12.24.$ ). When a plausible cause of speaker distraction was present, participants' fixation biases towards the distractor during disfluent utterances were reduced ( $\beta = -0.20, SE = 0.07, t = -2.9.$ ). The presence of the speaker-distraction was not found to have an effect on the tendency to fixate on the referent for fluent utterances ( $t = -0.9.$ ).

### Mouse-movements.

The proportion of mouse movements (distance traveled) towards either object was computed for each 20ms bin from referent onset to 2000ms post-onset. Figures 4 and 5 and show the proportion of mouse movements towards either object in fluent and disfluent conditions respectively. Model analysis of the window from referent-onset to 800ms post-onset patterned with the eye-tracking data. When presented with a fluent utterance, participants' movements to the distractor decreased over time in favour of the referent ( $\beta = -0.46, SE = 0.02, t = -19.49.$ ), whereas movements to the distractor over the referent increased over time when the delivery was disfluent ( $\beta = 0.60, SE = 0.03, t = 17.72.$ ). When

a plausible cause of speaker distraction was present, participants bias towards moving to the distractor during disfluent utterances was reduced ( $\beta = -0.35$ ,  $SE = 0.05$ ,  $t = -7.25$ ).

### Discussion.

Evidence of listeners' pragmatic judgments about a speaker's honesty based on manner of delivery emerged at an early stage in utterance processing, with eye- and mouse-tracking data showing biases towards referent and distractor for fluent and disfluent utterances respectively during the initial stage of the comprehension process. In this respect, the present study successfully replicates the effect found by Loy et al. [2016], showing that an utterance's delivery influences listeners' pragmatic inferences about speaker's intentions to be (dis)honest. Additionally, the results here show that this effect is robust against the presentation of speech in a noisy environment, in which there are potential distractions for the listener.

The interest in this study, however, is in how listeners responded when presented with two competing explanations for the cause of speaker-disfluency. When an alternative explanation for speaker-disfluency was available to listeners, although the disfluency-dishonesty bias persisted in the final interpretation of an utterance (as evidenced by the eventual decision of which object to click), the picture during the initial stages of comprehension is more complex. For disfluent utterances, the increase in eye- and mouse-movements to the distractor for disfluent conditions—indicating the emergent bias towards interpreting the utterance as dishonest—was initially attenuated when a possible cause of speaker-distraction was present. The availability of an external cause of the speaker disfluency resulted in a tendency to initially fixate and move toward the referent, only later fixating on and selecting the distractor.

One account of these findings would suggest that participants made early inferences about the contextual causes of disfluencies, which were eventually overridden by a dishonesty bias for disfluent utterances. By this account, listeners appear able to dynamically decide upon the most likely explanation of disfluency, and, as the speech unfolds, make attributions about why a speaker was disfluent. This equates to listeners engaging in a form of speaker-modelling, in which judgments are made about a speaker's perspective and thought-processes. These judgments are then used to constrain the interpretation of the speaker's meaning and intentions. The explanation of these results which suggests that listeners' disfluency-deception bias is reduced as a result of inferentially attributing the disfluency to the speaker being distracted by the car-horn, implies that listeners access information about a speaker's perspective (i.e. that they are distracted) during the early

stages of reference-comprehension.

However, it remains possible to explain these results in terms of direct associations of disfluencies and properties of language use. A probabilistic model of associations of disfluency and e.g. deception, discourse-new or unfamiliar objects, would require a sensitivity to proximate environmental interference such that the weight given to other disfluency-associations is reduced. The distinction here is that disfluency effects on comprehension may still be a learned response to general language use, rather than the result of listeners' attributing beliefs and intentions to particular speakers. Such an explanation may, however, still require listeners to make some form of judgment about what contextual information justifies a reduction in a disfluency's influence—would the sensitivity listeners display to contextual information when processing disfluency discern between the car-horn used here, which was situated within the context of the speaker's environment, and the same car horn if it occurred in the street outside the lab and was audible to the listener but external to the experiment? It is hard to see how a judgment about what constitutes plausible speaker-distraction differs from an inference about the cause of a speaker's disfluency.

The study here found that utterance-local contextual information influenced listeners' moment-to-moment pragmatic inferences about disfluencies. Whilst being eventually overridden by listeners' bias towards interpreting disfluent utterances as dishonest, the availability of an alternative, external cause of speaker-disfluency (being distracted) resulted in an initial attenuation of this bias. Whilst this contextual effect could be explained as a prelearned expectation about disfluency frequently occurring after a potentially distracting noise, it is not clear how listeners judge what might count as potentially distracting for a speaker. Alternatively, these findings could suggest that participants made early attributions about the potential causes of a disfluency, implying that listeners take into account a speaker's perspective during the initial stages of comprehension.

## References

- Alibali, M. W., Heath, D. C., and Myers, H. J. (2001). Effects of Visibility between Speaker and Listener on Gesture Production: Some Gestures Are Meant to Be Seen. *Journal of Memory and Language*, 44(2):169–188.
- Arnold, J. E., Fagnano, M., and Tanenhaus, M. K. (2003). Disfluencies Signal Theee, Um, New Information. *Journal of Psycholinguistic Research*, 32(1):25–36.
- Arnold, J. E., Kam, C. L. H., and Tanenhaus, M. K. (2007). If you say thee uh you are describing something hard: The on-line attribution of disfluency during reference comprehension. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 33(5):914–930.
- Arnold, J. E., Tanenhaus, M. K., Altmann, R. J., and Fagnano, M. (2004). The Old and Thee, uh, New. *Psychological Science*, 15(9):578–582.
- Barr, D. and Seyfeddinipur, M. (2010). The role of fillers in listener attributions for speaker disfluency. *Language and cognitive processes*, 25(4):441–455.
- Barr, D. J. (2001). Trouble in mind: Paralinguistic indices of effort and uncertainty in communication. In Cavé, C., Guaïtella, I., and Santi, S., editors, *Oralité and gestualité: Interactions et comportements multimodaux dans la communication*, pages 597–600. Paris: L’Harmattan.
- Barr, D. J. (2004). Establishing conventional communication systems: Is common knowledge necessary? *Cognitive Science*, 28(6):937–962.
- Barr, D. J. (2008). Analyzing ‘visual world’ eyetracking data using multilevel logistic regression. *Journal of Memory and Language*, 59(4):457–474.
- Barr, D. J. (2014). Perspective Taking and its Impostors in Language Use: Four Patterns of Deception. *The Oxford Handbook of Language and Social Psychology*, pages 98–110.
- Bates, D., Mächler, M., Bolker, B., and Walker, S. (2015). Fitting Linear Mixed-Effects Models Using lme4. *Journal of Statistical Software*, 67(1):1–48.
- Beattie, G. W. (1979). Planning units in spontaneous speech: Some evidence from hesitation in speech and speaker gaze direction in conversation. *Linguistics*, 17(1-2):61–78.

- Bond, C. F. and DePaulo, B. M. (2006). Accuracy of deception judgments. *Personality and social psychology review : an official journal of the Society for Personality and Social Psychology, Inc*, 10(3):214–234.
- Bortfeld, H., Leon, S. D., Bloom, J. E., Schober, M. F., and Brennan, S. E. (2001). Disfluency Rates in Conversation: Effects of Age, Relationship, Topic, Role, and Gender. *Language and Speech*, 44(2):123–147.
- Brennan, S. and Williams, M. (1995). The Feeling of Another’s Knowing: Prosody and Filled Pauses as Cues to Listeners about the Metacognitive States of Speakers. *Journal of Memory and Language*, 34(3):383–398.
- Clark, H. H. (1996). Using language. 1996. *Cambridge University Press: Cambridge*, 952:274–296.
- Clark, H. H. and Fox Tree, J. E. (2002). Using uh and um in spontaneous speaking. *Cognition*, 84(1):73–111.
- Cook, S. W., Jaeger, T. F., and Tanenhaus, M. K. (2009). Producing Less Preferred Structures: More Gestures, Less Fluency. *31st Annual Conference of the Cognitive Science Society*, pages 62–67.
- Corley, M., MacGregor, L. J., and Donaldson, D. I. (2007). It’s the way that you, er, say it: Hesitations in speech affect language comprehension. *Cognition*, 105(3):658–668.
- Eberhard, K. M., Spivey-Knowlton, M. J., Sedivy, J. C., and Tanenhaus, M. K. (1995). Eye movements as a window into real-time spoken language comprehension in natural contexts. *Journal of Psycholinguistic Research*, 24(6):409–436.
- Fernald, A. and Mazzie, C. (1991). Prosody and focus in speech to infants and adults. *Developmental Psychology*, 27(2):209–221.
- Fox Tree, J. E. (1995). The Effects of False Starts and Repetitions on the Processing of Subsequent Words in Spontaneous Speech. *Journal of Memory and Language*, 34(6):709–738.
- Frazier, L., Carlson, K., and Clifton, C. (2006). Prosodic phrasing is central to language comprehension. *Trends in Cognitive Sciences*, 10(6):244–249.



- Heller, D., Arnold, J. E., Klein, N., and Tanenhaus, M. K. (2015). Inferring Difficulty: Flexibility in the Real-time Processing of Disfluency. *Language and Speech*, 58(2):190–203.
- Krahmer, E. and Swerts, M. (2005). How Children and Adults Produce and Perceive Uncertainty in Audiovisual Speech. *Language and Speech*, 48(1):29–53.
- Kronmüller, E. and Barr, D. J. (2007). Perspective-free pragmatics: Broken precedents and the recovery-from-preemption hypothesis. *Journal of Memory and Language*, 56(3):436–455.
- Lin, S., Keysar, B., and Epley, N. (2010). Reflexively mindblind: Using theory of mind to interpret behavior requires effortful attention. *Journal of Experimental Social Psychology*, 46(3):551–556.
- Loy, J. E., Rohde, H., and Corley, M. (2016). Effects of Disfluency in Online Interpretation of Deception. *Cognitive Science*.
- Mathôt, S., Schreij, D., and Theeuwes, J. (2012). OpenSesame: An open-source, graphical experiment builder for the social sciences. *Behavior Research Methods*, 44(2):314–324.
- Pickering, M. J. and Garrod, S. (2004). Toward a mechanistic psychology of dialogue. *The Behavioral and brain sciences*, 27(2):169–190; discussion 190–226.
- R Core Team (2016). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria.
- Schachter, S., Christenfeld, N., Ravina, B., and Bilous, F. (1991). Speech disfluency and the structure of knowledge. *Journal of Personality and Social Psychology*, 60(3):362–367.
- Smith, V. L. and Clark, H. H. (1993). On the Course of Answering Questions. *Journal of Memory and Language*, 32(1):25–38.
- Snodgrass, J. G. and Vanderwart, M. (1980). A standardized set of 260 pictures: Norms for name agreement, image agreement, familiarity, and visual complexity. *Journal of Experimental Psychology: Human Learning & Memory*, 6(2):174–215.
- Swerts, M. and Krahmer, E. (2005). Audiovisual prosody and feeling of knowing. *Journal of Memory and Language*, 53(1):81–94.
- Zuckerman, M., Koestner, R., and Driver, R. (1981). Beliefs about cues associated with deception. *Journal of Nonverbal Behavior*, 6(2):105–114.

Filler type	Manipulation	No. of Utterances	Example
Fluent	None	20	The treasure is behind the <referent>.
Disfluent	Prolongation	3	The treasure is behind <i>thee</i> ... <referent>.
	Repetition	4	The treasure is behind <i>the- the</i> <referent>.
	Filled pause (utterance-initial)	3	<i>Umm..</i> The treasure is behind the <referent>.
Other	Discourse marker	5	<i>Okay,</i> the treasure is behind the <referent>.
	Modal	3	The treasure <i>could be</i> behind the <referent>.
	Combination	2	<i>Right,</i> the treasure <i>might be</i> behind the <referent>.

Table 1

*Disfluencies and discourse manipulations in filler items.*

Noise	No. of Utterances
Vehicle horns (various)	7
Sirens (various)	3
Vehicles revving (various)	2
Car-stereo	2
Bicycle-bell	1
Bus doors opening	1
Footsteps	1
Loose drain cover	1
Man shouting	1
Dog barking	1

Table 2

*Plausible causes of speaker distraction in filler items.*

Did you find that these background noises were distracting for you as a listener?	58%
Did you think that these background noises affected the speaker?	50%
Did you think that the background noises, in some cases, might have caused the speaker to be disfluent (to say 'um' or 'uh')?	50%

Table 3

*Percentage of participants who responded yes/possibly/probably.*

	<b>Delivery</b>	Disfluent	Disfluent	Fluent	Fluent
	<b>Speaker-distraction</b>	Absent	Present	Absent	Present
Distractor		62%	62%	23%	24%
Referent		38%	38%	77%	76%

Table 4

*Mouse clicks on each object by condition.*

Effect	$\beta$	SE	t-value
(Intercept)	-0.60783	0.06074	-10.007
Time	0.33195	0.03640	9.118
Object	0.25363	0.03376	7.512
Delivery	0.09023	0.03384	2.666
Distraction	-0.11175	0.03377	-3.309
Time:Object	-0.62502	0.03376	-18.512
Time:Delivery	-0.21521	0.03384	-6.359
Object:Delivery	-0.32129	0.04785	-6.715
Time:Distraction	0.05362	0.03377	1.588
Object:Distraction	0.11787	0.04775	2.469
Delivery:Distraction	0.11588	0.04808	2.410
Time:Object:Delivery	0.58559	0.04785	12.238
Time:Object:Distraction	-0.04478	0.04775	-0.938
Time:Delivery:Distraction	-0.02308	0.04808	-0.480
Object:Delivery:Distraction	-0.02619	0.06797	-0.385
Time:Object:Delivery:Distraction	-0.19942	0.06797	-2.934
<b>Random effects</b>			
	Var.	S.D	Corr.
By-Subject			
(Intercept)	0.070	0.264	
Time	0.013	0.116	-0.720
By-Item			
(Intercept)	0.004	0.066	
Time	0.003	0.063	-0.940
Residuals			
	0.636	0.798	

Table 5

Effect	$\beta$	SE	t-value
(Intercept)	-1.19815	0.03593	-33.35
Time	0.57152	0.03557	16.07
Object	0.17762	0.02370	.49
Delivery	0.22830	0.02400	9.51
Distraction	0.02378	0.02397	0.99
Time:Object	-0.46372	0.02379	-19.49
Time:Delivery	-0.28300	0.02407	-11.76
Object:Delivery	-0.31433	0.03392	-9.27
Time:Distraction	-0.08203	0.02400	-3.42
Object:Distraction	-0.04261	0.03388	-1.26
Delivery:Distraction	-0.04890	0.03411	-1.43
Time:Object:Delivery	0.60288	0.03402	17.72
Time:Object:Distraction	0.13227	0.03391	3.90
Time:Delivery:Distraction	0.21064	0.03411	6.18
Object:Delivery:Distraction	-0.09102	0.04821	1.89
Time:Object:Delivery:Distraction	-0.34941	0.04821	-7.25
<b>Random effects</b>			
	Var.	S.D	Corr.
By-Subject			
(Intercept)	0.020	0.141	
Time	0.016	0.128	0.110
By-Item			
(Intercept)	0.004	0.060	
Time	0.006	0.078	-0.790
Residuals			
	0.346	0.589	


Table 6




"The treasure is behind thee, um --- bell"

*Figure 1.* Screenshot of video presented to participants as an example of speaker producing utterances in a noisy environment. Screenshot shows an example of speaker producing a disfluent and dishonest utterance.







*Figure 2.* Mean proportion of fixations to either object (referent and distractor) for fluent utterances split by presence of speaker distraction, calculated out of the total sum of fixations for each 20ms time bin from referent-onset to 2000ms post-onset. Shaded areas represent  $\pm 1$  standard error of the mean. Time bins in which a bias towards either the referent or distractor was evident are highlighted with respect to condition and direction of bias (assessed by a mixed effect model estimating the empirical logit of fixations to either object in a given condition, with main effect of object and by-subject and by-item random intercepts and random slopes of object).



*Figure 3.* Mean proportion of fixations to either object (referent and distractor) for disfluent utterances split by presence of speaker distraction, calculated out of the total sum of fixations for each 20ms time bin from referent-onset to 2000ms post-onset. Shaded areas represent  $\pm 1$  standard error of the mean. Time bins in which a bias towards either the referent or distractor was evident are highlighted with respect to condition and direction of bias (assessed by a mixed effect model estimating the empirical logit of fixations to either object in a given condition, with main effect of object and by-subject and by-item random intercepts and random slopes of object).



*Figure 4.* Mean proportion of cumulative distance traveled toward each object (referent or distractor) in fluent conditions split by presence of speaker distraction, from referent onset to 2000ms post-onset. Proportions calculated out of total cumulative distance moved the mouse from referent-onset until that time bin. Shaded areas represent  $\pm 1$  standard error of the mean.



*Figure 5.* Mean proportion of cumulative distance traveled toward each object (referent or distractor) in disfluent conditions split by presence of speaker distraction, from referent onset to 2000ms post-onset. Proportions calculated out of total cumulative distance moved the mouse from referent-onset until that time bin. Shaded areas represent  $\pm 1$  standard error of the mean.